

# Deep learning in Finance

Damien Challet  
[damien.challet@centralesupelec.fr](mailto:damien.challet@centralesupelec.fr)

22nd October 2024

# Course philosophy

*DL from the point of view of Finance*

# Course philosophy

- How/when not to use DL
- Methodology:
  - good filtering?
  - pre-knowledge?
- how to talk to DL
  - inputs
  - architecture
  - check default options
  - result interpretation ↔ filtering
- Mark: 100% from TPs.

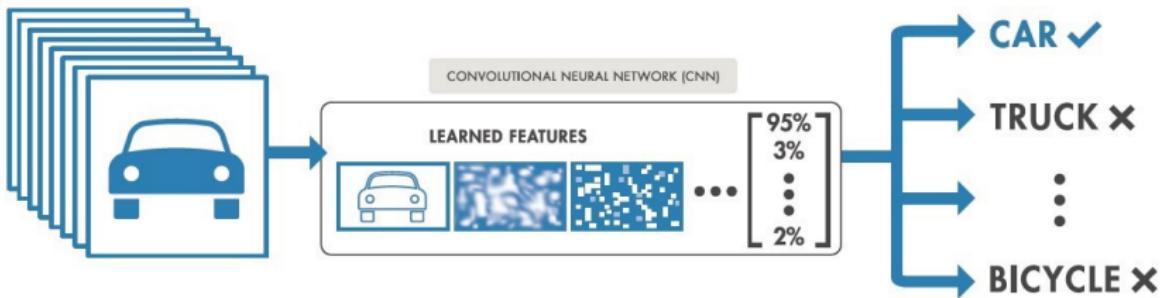
# Success of Deep Learning

- Image
- Sound
- Text
- Image to text, text to image, sound to text, ...

State of AI annual report (206 pages) [link]

- Finance?

# DL: image classification



[source]

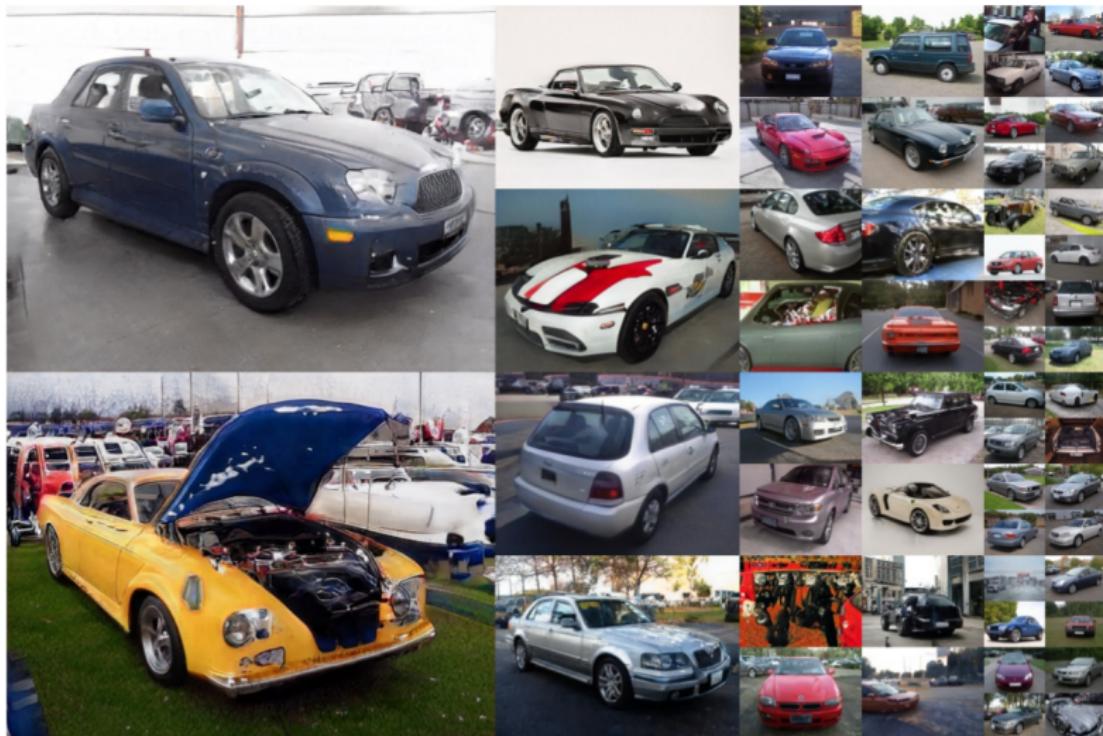
# DL: image generation

[source], see also [this person does not exist]



# DL: image generation

[source]



# DL: text to image (2021)

Google red classic car open bonnet

All Images Maps Videos Shopping More Tools

stock photo 123rf convertible voiture aston martin american

Red Classic Vintage Car Stock Photo ...  
123rf.com

Red Vintage Car With Open Hood Stock ...  
alamy.com

Classic car with open hood stock image ...  
dreamstime.com

Red Classic Automobile With Open Hood  
iStock

Vintage British Sports Car Hood Up  
alamy

Red classic car with open hood and ...  
alamy

# DL: text to image (2024)

Stable diffusion (automatic111): 'Beach, morning, seagulls, pizza'



## The art of the prompt

glowing galaxy cracks, A cute, glowing version of SpongeBob SquarePants, designed with irresistible, large eyes. His body appears to be cracking like a molten volcano, with cracks glowing in bright, galaxy-like colors—shades of deep purple, blue, and bright orange. The glow emanates from within the cracks, giving an ethereal and mesmerizing effect, as if SpongeBob is infused with cosmic energy. The background is dark, making the glow and molten effect stand out, and small, floating star-like particles surround him, enhancing the galaxy and volcanic theme. The scene blends cuteness with a cosmic, otherworldly vibe

# The art of the prompt

[civitai.com]

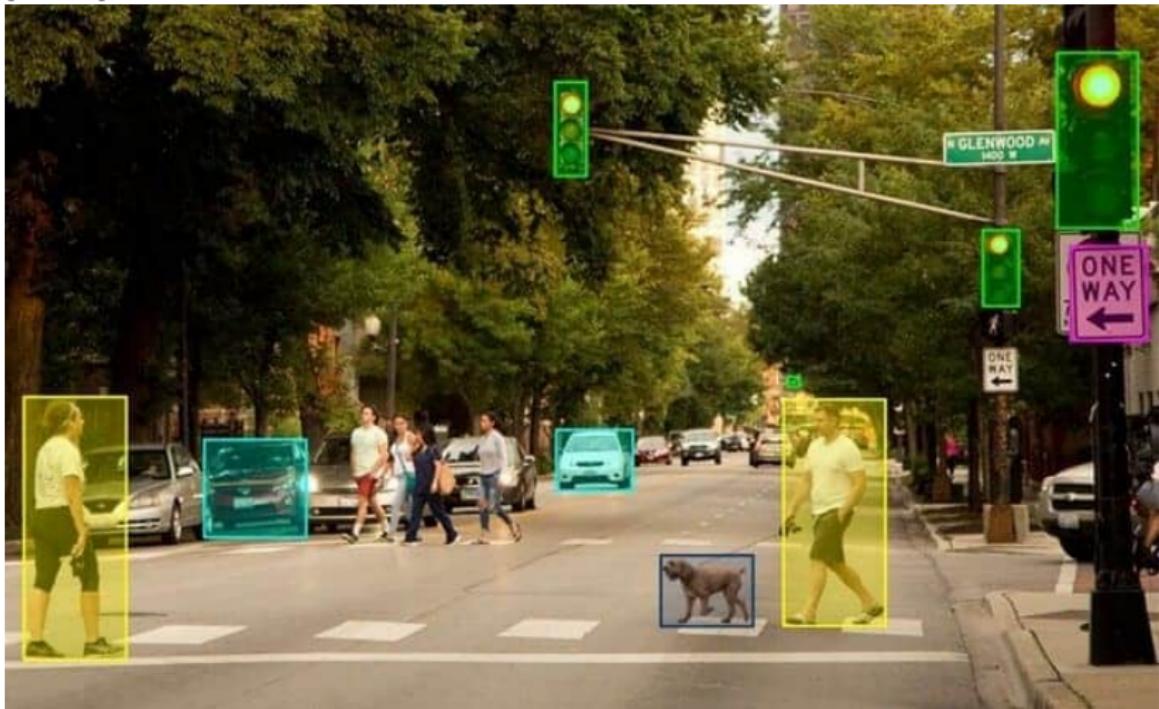


and of the model



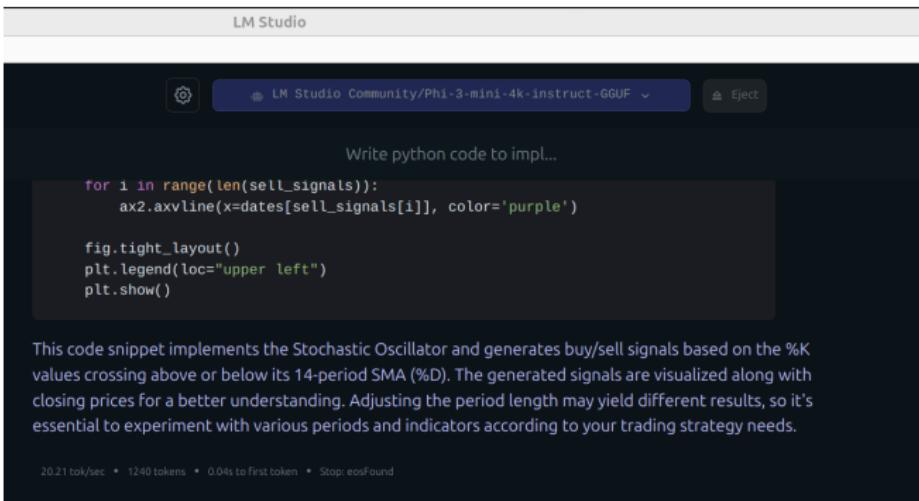
# Image annotation

[source]



# LLMs on your laptop / (mobile phone)

LMstudio [link]



The screenshot shows the LM Studio application window. At the top, there's a toolbar with icons for file operations and a dropdown menu showing "LM Studio Community/Phi-3-mini-4k-instruct-GGUF". Below the toolbar, a status bar displays "Eject". The main area has a dark background with a light gray input field containing placeholder text "Write python code to impl...". A code completion suggestion is shown in a semi-transparent box: 

```
for i in range(len(sell_signals)):  
    ax2.axvline(x=dates[sell_signals[i]], color='purple')  
  
fig.tight_layout()  
plt.legend(loc="upper left")  
plt.show()
```

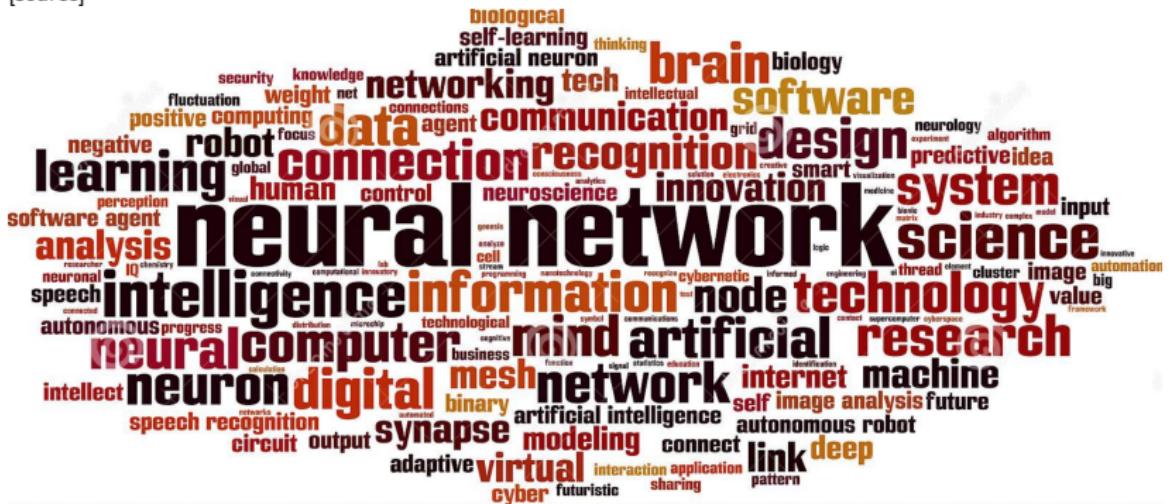
 Below this, a descriptive text block reads: "This code snippet implements the Stochastic Oscillator and generates buy/sell signals based on the %K values crossing above or below its 14-period SMA (%D). The generated signals are visualized along with closing prices for a better understanding. Adjusting the period length may yield different results, so it's essential to experiment with various periods and indicators according to your trading strategy needs." At the bottom of the window, performance metrics are displayed: "20.21 tok/sec \* 1240 tokens \* 0.04s to first token \* Stop: eosFound".

# Course plan

- |                                |                     |
|--------------------------------|---------------------|
| 1. Estimation/calibration      | MLP/CNN             |
| 2. Option pricing              | model vs DL         |
| 3. Market state classification | auto-encoders       |
| 4. Timeseries generation       | GANs                |
| 5. Timeseries generation 2     | GANs                |
| 6. Prediction                  | RNN                 |
| 7. Prediction 2                | RNN/CNN/VAE/MLP/LTR |
| 8. Oral feedback on TPs        |                     |

Deep learning from the point of view of finance

[source]



© dreamstime.com

ID 187912661 © Boris15

Finance from the point of view of Data Science

# YAMLP

*Yet Another ML Project*

# Finance and Machine Learning

- Traditional finance is very linear:

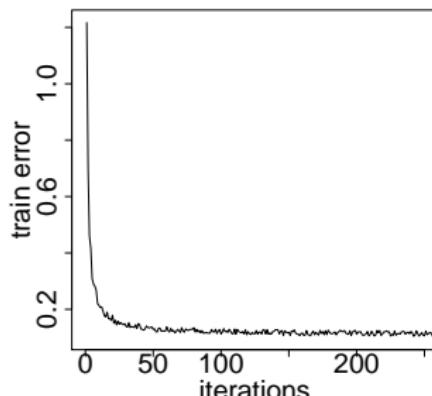
$$\text{CAPM: } \hat{r}_{i,t} = \alpha_i + \beta_i r_{market,t} + \eta_{i,t}$$

$$\text{Factors: } \hat{r}_{i,t} = \alpha_i + \sum_k w_{i,k} f_{i,k,t} + \eta_{i,t}$$

- ML: adds non-linearity
  - Support Vector Machines
  - Genetic programming
  - Random Forests
  - XGboost, LightGBM, ...
  - **DL**

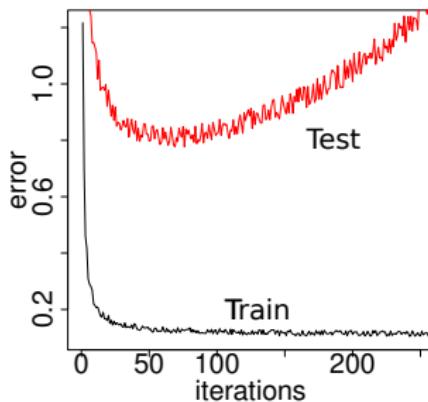
## YAMLP: minimal approach to success in ML

1. Split data in train / test sets (80/20)
2. Train ML model on train set until error stops decreasing
3. Test model on test set  
(~fine for large stationary synthetic data)



## YAMLP: over-fitting

1. Split data in train / test sets (80/20)
2. Train ML model until test error stops decreasing (early stopping)
3. Report test error



# YAMLP: a more subtle approach to success in ML

1. Split data in train / validation / test sets (60/20/20)
2. Train ML model until *validation* error stops decreasing
3. Test model

# YAMLP: rescaling

Types of rescaling:

- MinMax rescaler

$$\tilde{x}_i = \frac{x_i - \min_k x_k}{\max_j x_j - \min_k x_k} \in [0, 1]$$

- Z-score

$$\tilde{x}_i = \frac{x_i - \text{mean}(\mathbf{x})}{\text{std}(\mathbf{x})}$$

## YAMLP successful: sharp stationary truth

$X$  input,  $y$  output

$$y = F(X)$$

1. Sharp  $X$  and  $y$
2. Stationary world:  $F$  constant

## Finance is NOT YAMLP (Yet Another ML Project)

## Specific problems

- Heavy tailed data rescaling?
  - Non-stationarity distribution shift
  - Long memory ??
  - Good-enough price predictions are *really hard* [link]

# Problem with rescaling: heavy-tailed data

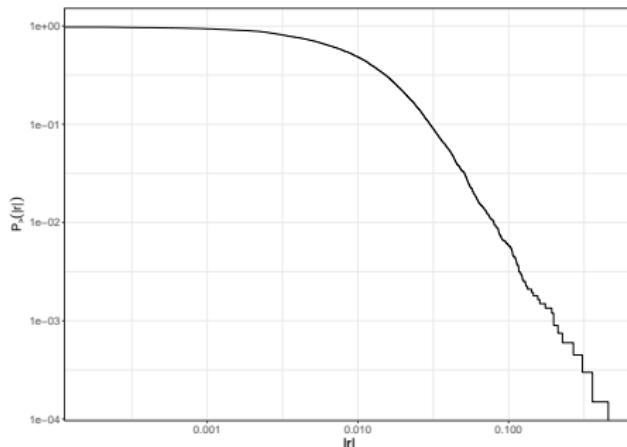
- Price return

$$R_t = p_t / p_{t-1} - 1$$

- Price log return

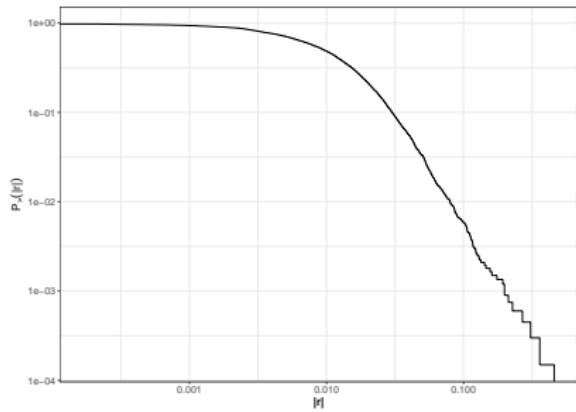
$$r_t = \log p_t - \log p_{t-1} = \log(1 + R_t)$$

- Heavy tail:  $P(|\text{return}| > r) \propto |r|^{-3}$



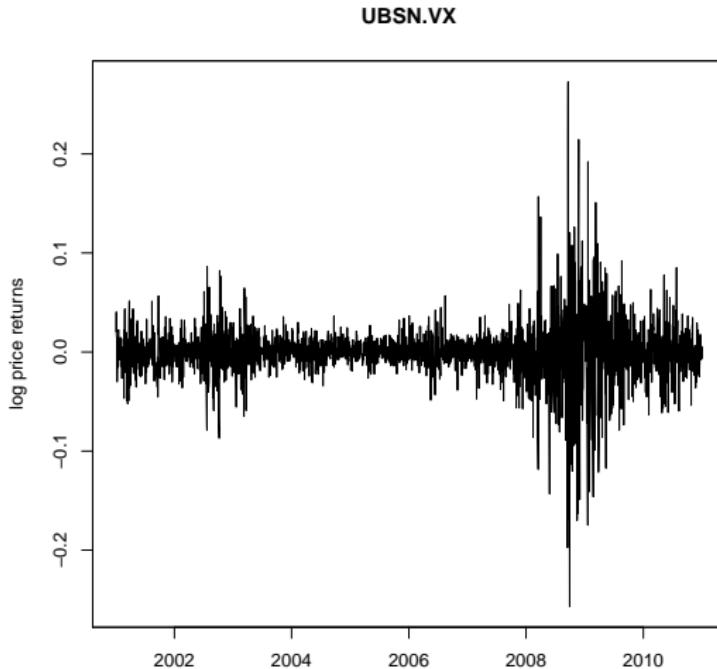
# Problem with rescaling: heavy-tailed data

- MinMax:
  - unbounded values
  - max and min are not robust
- Z-score:
  - Usually  $E(|r|^3) \rightarrow \infty$
  - StdDev estimator is too noisy
  - Average estimator is too noisy



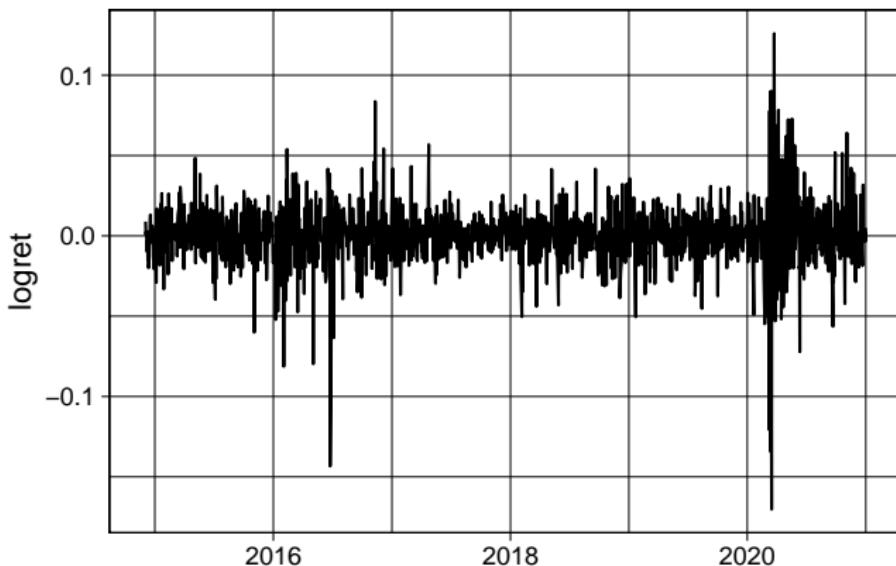
→ naive input rescaling may not be a wise strategy

## Problem with rescaling: non-stationarity



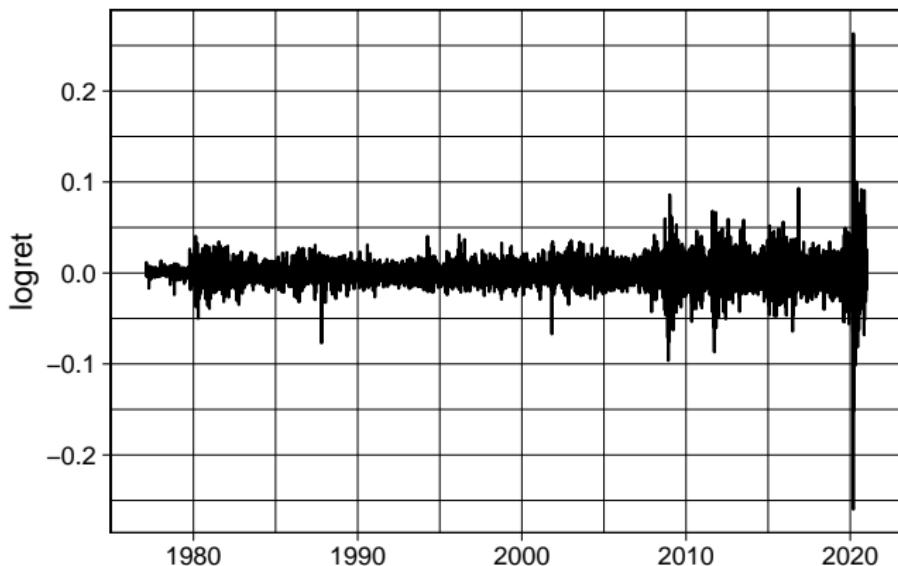
## Non-stationarity is non-stationary

UBS



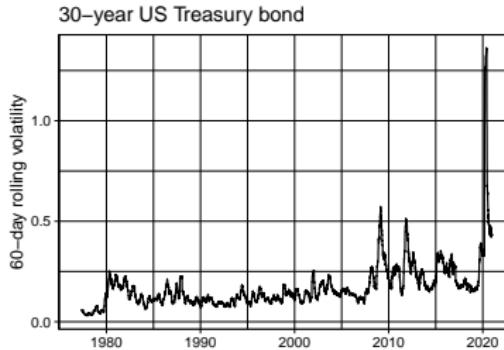
All asset types are non-stationary and heavy-tailed

30-year US Treasury bond



# Problem with full data rescaling: information leakage

- Many ML blogs make this mistake
- Rescaling must be done within each set, not globally
- ANY information from the future makes prediction significantly biased.
- Accounting for future new scale maxima uses the future



## Problem with single train window: non-stationarity

- Many papers/blogs only train once
- Dependence structure of financial data is non-stationary
  - train/validation/test once is not enough
  - test length = real-life trading: as short as possible.
- Problem: DL may be prohibitively expensive
- Solutions:
  - stationarize ML model: very long time series
  - warm start
  - predict distribution shifts

## Problem with simple NNs: time-series, long memory

- Problem: Many financial quantities have a long memory
  - volatility
  - volume
  - order flow (e.g. market order sign)
- Solution: use adequate model / NN architecture
  - Temporal Convolution Nets (TCNs) (Wavenets)
  - LSTM / GRU
  - State-space models (Mamba)

## Problem with predictions:

- Many quantities are hard to predict.
- Many ML blogs/articles find  $\hat{p}_{t+1} = p_t \iff E(r_t) = 0$ : no prediction power
- Any small information leakage from the future changes much the results.
  - scale from the future
  - train/test split not causal
  - over-optimizing ML (data snooping)

## Problem with predictions: time series are VERY noisy

- Noise filtering is essential before using a matrix of price returns (e.g.)
- Real information content of a single price return is very small: 1-3 bits, not 64.
- *Purgamentum init, purgamentum exit*

# Problem with predictions: literature very noisy

The screenshot shows a Google Scholar search results page. The search query is "llm "price prediction" trading". There are 273 results found in 0.10 seconds. The results are listed below, each with a title, abstract, and download links.

**Articles** About 273 results (0.10 sec)

**Any time**  
Since 2024  
Since 2023  
Since 2020  
Custom range...

**Sort by relevance**  
Sort by date

**Any type**  
Review articles

include patents  
 include citations

Create alert

**Carbon Price Forecasting with LLM-Based Refinement and Transfer-Learning** [\[PDF\]](#) [github.io](#)  
H Jiang, Y Ding, R Chen, C Fan - International Conference on Artificial ..., 2024 - Springer  
... feature selection and carbon futures **price prediction** in EU ETS market... However, the research on the **price prediction** of China's ... We collect from each CEA market the closing price, **trade** ...  
☆ Save 99 Cite Related articles Import into BibTeX

**Application and Modeling of LLM in Quantitative Trading Using Deep Learning Strategies**  
T Pan, J Yu, L Zheng, Y Li - Biologically Inspired Cognitive Architectures ..., 2023 - Springer  
... This project trained a stock **price prediction** model using long short-term memory (LSTM) methods. Then, the backtest model was established with the classic double moving average ...  
☆ Save 99 Cite Related articles Import into BibTeX

**Stock Price Trend Prediction using Emotion Analysis of Financial Headlines with Distilled LLM Model** [\[PDF\]](#) [acm.org](#)  
R Bhat, B Jain - Proceedings of the 17th International Conference on ..., 2024 - dl.acm.org  
... The stock **price prediction** domain has undergone significant research and debate, with ... stock price during the **trading** window or when the market opens for **trading**. Researchers have ...  
☆ Save 99 Cite Related articles Import into BibTeX

**Large Language Model Agent in Financial Trading: A Survey** [\[PDF\]](#) [arxiv.org](#)  
H Ding, Y Li, J Wang, H Chen - arXiv preprint arXiv:2408.06361, 2024 - arxiv.org  
... media data for stock **price prediction** such as [14, 40, 58]. However, SEP [19] is the only work we've reviewed that incorporates real-time social media data in **LLM trading** agent. In SEP, ...  
☆ Save 99 Cite Cited by 2 Related articles All 4 versions Import into BibTeX

**A Reflective LLM-based Agent to Guide Zero-shot Cryptocurrency Trading** [\[PDF\]](#) [arxiv.org](#)  
oct. 21 13:35

# DL for hard problems in Finance

- Exact nonlinear results are rare approx. solution
- Monte-Carlo simulations are expensive approx. solution
- Calibration is hard and expensive DL surrogates
- Prediction is harder hidden relationships
- Many equations are impossible to solve PINNs

# Artificial Neural Networks: a short introduction

# Artificial Neural Networks: a short introduction

- NN = collection of *connected* neurons
- Artificial neuron:  
weighted sum of inputs + non-linear *activation*

$$O = \phi \left( \sum_{i=1}^N w_i x_i \right)$$

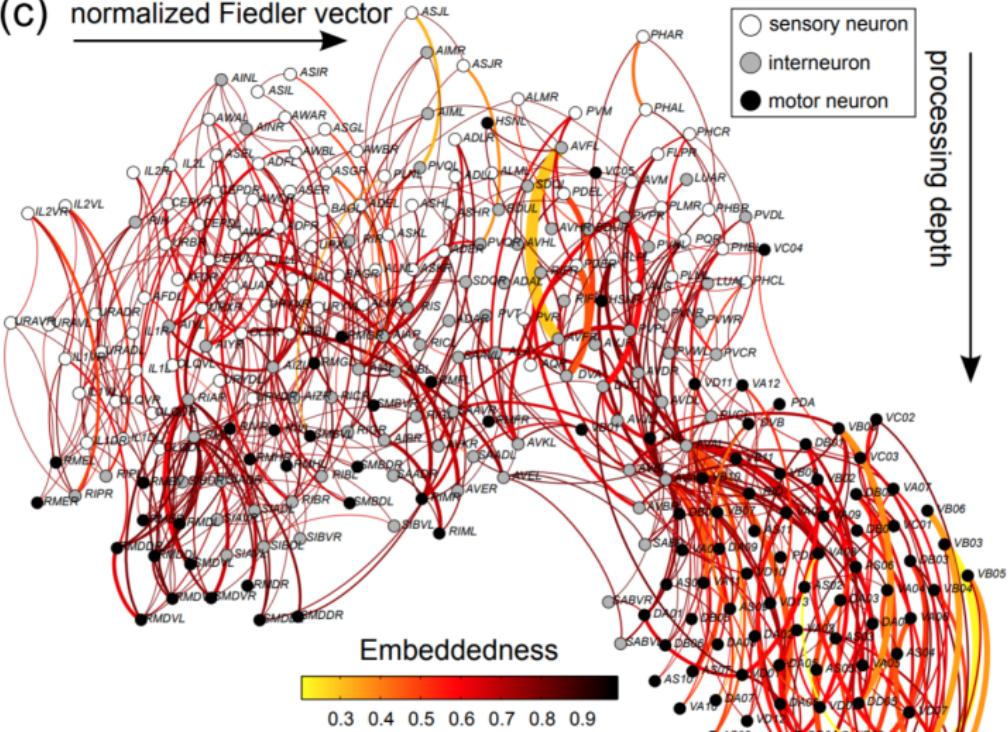
$\phi$  : non-linear function

- What happens if many non-linear functions interact?
  - complex system (simple elements, complex behaviour)
  - even more complex if time-dependent with feed-back
  - even more complex if connectivity is complex

## Real-life connectivity

C. Elegans [Schaub et al. (2013)]

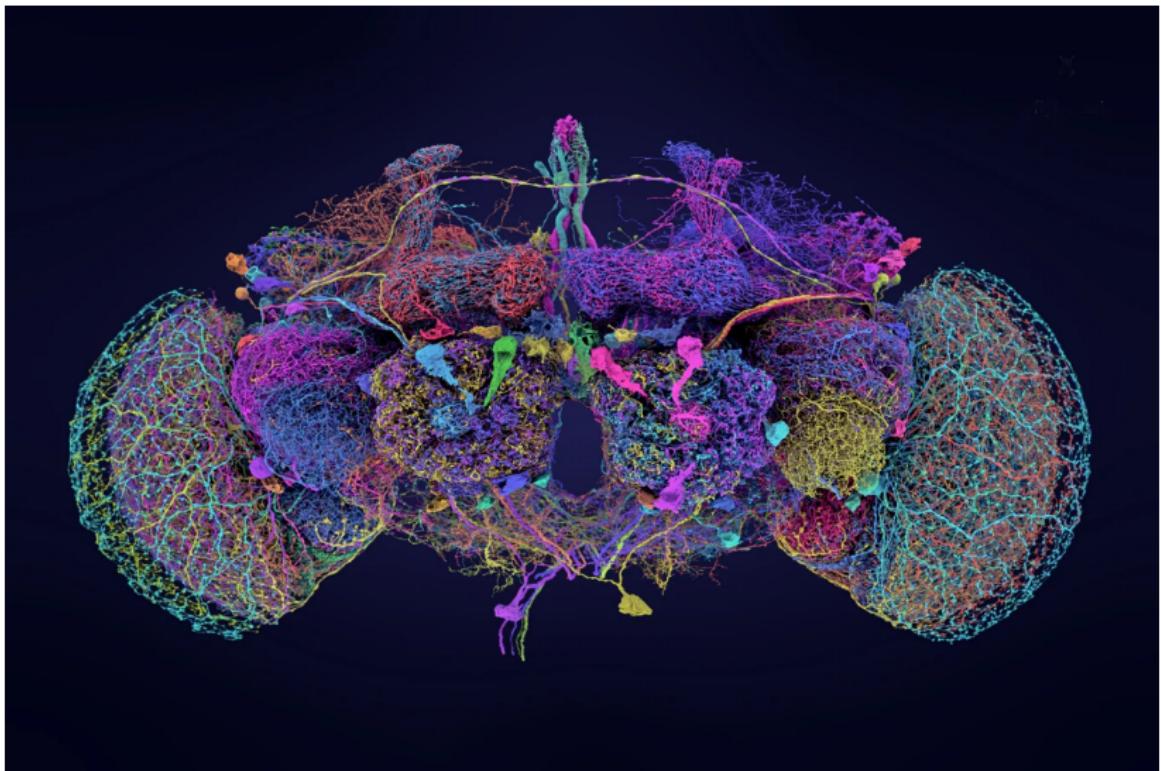
(c) normalized Fiedler vector



(d)

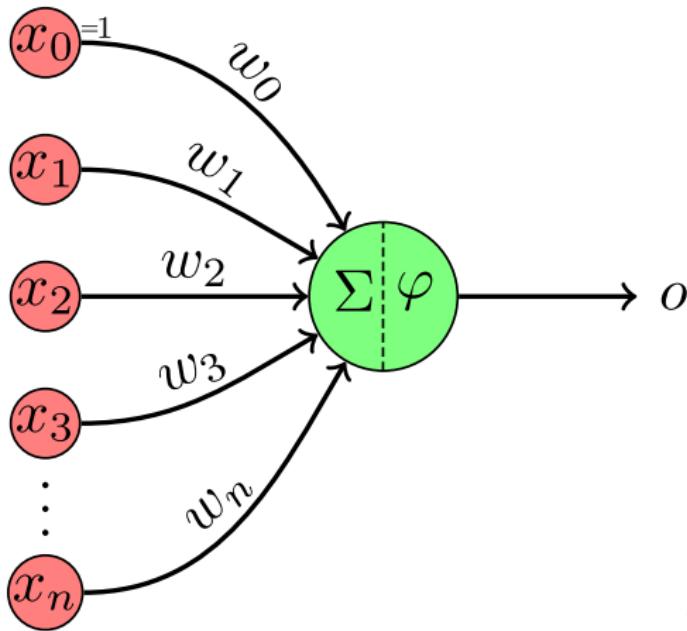
# Fruit fly (2024)

Dorkenwald *et al.* (2024)[link]: 139,255 neurons

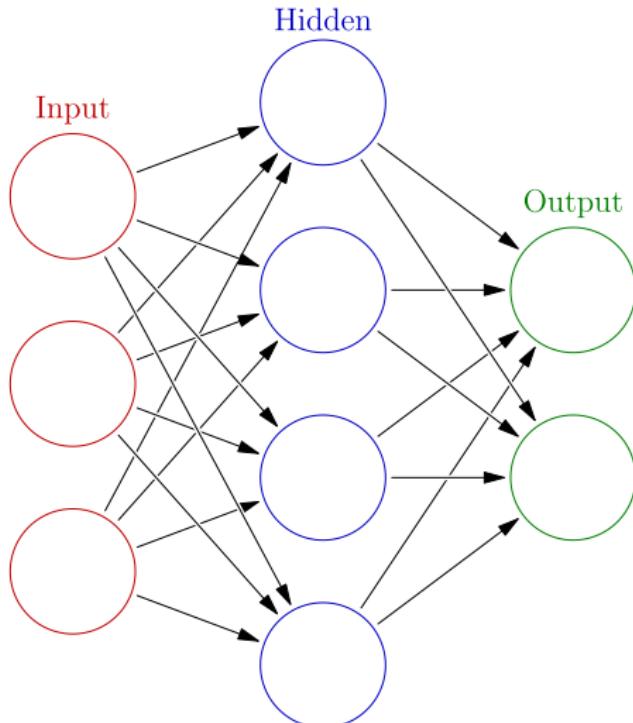


# ANNs: simplest perceptron

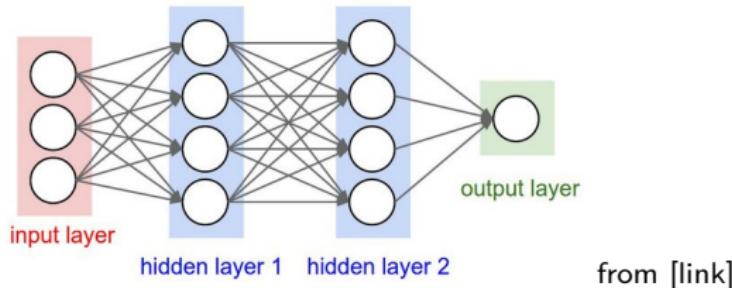
Single neuron



# Single-hidden layer ANN



# Why Multilayer ANNs (a.k.a. Deep)?



- Neuron budget:  $N$
- $N$  neurons in a single layer
- $K$  hidden layers with  $N/K$  neurons

$$\rightarrow \text{complexity } O\left[\left(\frac{N}{K}\right)^K\right] > O(N)$$

# Biological NNs

More complex architecture

- Shortcuts
- Recurrent
- Stateful neurons
- Time dimension
  - synchronisation?
  - functional programming
- Principle of least action (not least loss)

## Deep Neural Network: definition

at least 1 hidden layer + something to optimize

# What to optimize?

1. Least-squares
2. Non-linear fit
3. Estimation
4. Distance to complex equation solution
5. Risk minimization
6. Trading performance
7. Terminal wealth / integrated cost
8. Look-alike-ness time series
9. Minimise deviation from constraints

## Some notations

- $N$  samples of  $P$  features

$$X_{n,p}, \quad n = 1, \dots, N \quad p = 1, \dots, P$$

- Output  $Y$  of size  $K$

$$Y_n, \quad n = 1, \dots, N$$

$$(X_{n,1} \cdots X_{n,P}) \sim (Y_{n,1} \cdots Y_{n,K})$$

- ML: match lines of  $X$  with lines of  $Y$

$$X \sim Y$$

# Tasks

- Regression:

$$Y_n = F(X_n) \in \mathbb{R}^K$$

- Classification:  $S$  classes

$$Y_n = F(X_n) \in \{1, \dots, S\}^K$$

- Discrimination

$$F(X_n) \neq F(X_m), \quad n \neq m$$

- Ranking

$$F(X_n) > F(X_m), \quad n \neq m$$

# DL $\equiv$ error minimisation

Regression

$$\mathcal{L} = \sum_i \Phi [Y_i - F(X_i)]$$

- MSE ( $\rightarrow$  average)

$$\mathcal{L} = \sum_i [Y_i - F(X_i)]^2$$

- MAE ( $\rightarrow$  median)

$$\mathcal{L} = \sum_i |Y_i - F(X_i)|$$

- Robust loss (Huber)

$$\Phi(x) = \begin{cases} x & |x| < c \\ K & |x| \geq c \end{cases}$$

# How to optimise something with an ANN?

Backpropagation: only weights matters

Nice write-up [link]

- Feed forward:

Input → ANN → output

- error:

$$\mathcal{L} = E(|\text{output} - \text{expected output}| \mid \mathbf{w})$$

- backpropagation: weight update, layer  $i$

$$w'_i = w_i - \eta \frac{\partial \mathcal{L}}{\partial w_i}$$

$\eta$  : learning rate

- computing  $\frac{\partial \mathcal{L}}{\partial w_i} = \prod_{j \geq i} \frac{\partial \mathcal{L}}{\partial w_j}$  : chain rule

# Do biological networks backpropagate?

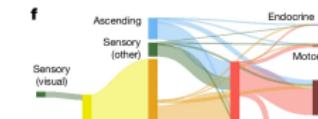
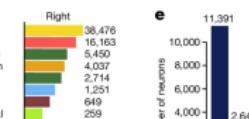
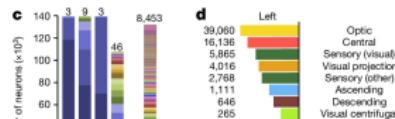
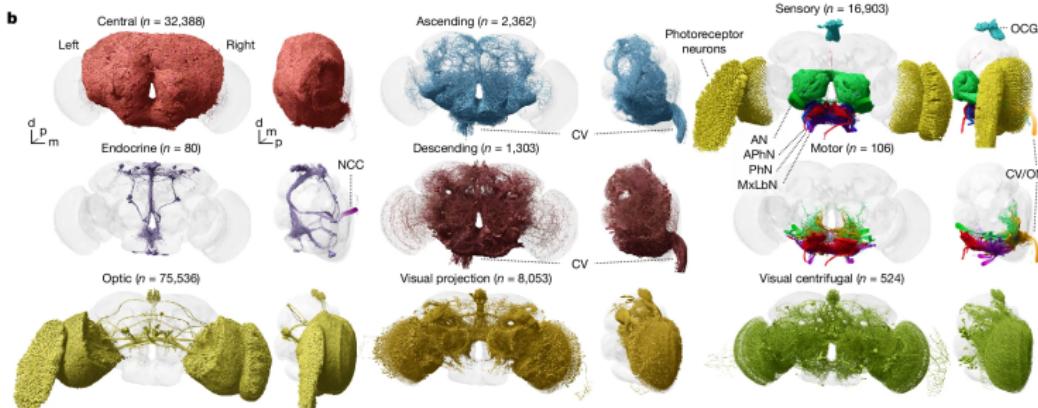
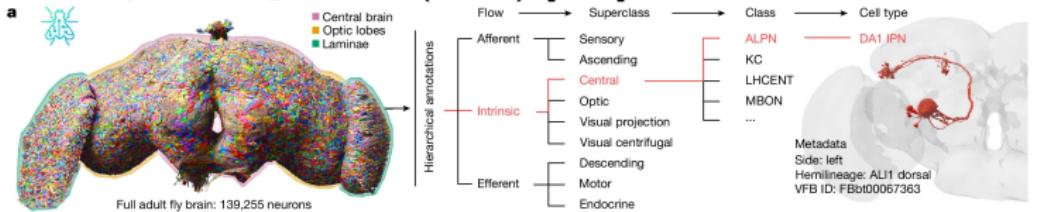
Alternative: predictive configuration/coding [link]

- Activation strength matter
  - Find better weights that yield better activation strengths
1. → update neuron activation: control variables
  2. → update weights: parameters

Song et al. (2023) [link]

# Backward-forward propagation

Fruit fly: Schlegel et al (2024) [link]



## Side note: alternative ways of learning

Bak and Chialvo: learning from mistakes (2001) [[link](#)]

- Classification
- Activation: neuron with  $\max w_i$  is activated
- Learning only from mistakes
  - reduce weights of activated neurons
  - quick re-learning

# DL success: learn representations

Convolutional networks [source]



- The deeper, the more and more complex the features
- Interactive visualization [[link](#)]

# DL programming: crash course

1. No installation: use Google Colab [[link](#)]  
free **fast** CPU / GPU / TPU  
(runtime->GPU)
2. Nvidia: use a fully set up Docker container from Nvidia  
AMD and Mac: install tensorflow
3. Install Tensorflow 2.x  
`conda install -c anaconda tensorflow`  
preferably in a virtual environment (`conda create -n DLF`)

# My first ANN: single-layer NN

Input:  $P = 8$ , output  $K = 1$

```
from tensorflow import keras
from tensorflow.keras.models import Sequential
from tensorflow.keras.layers import Dense

model = Sequential()

model.add(Dense(12, input_dim=8, activation='relu'))
model.add(Dense(1, activation='sigmoid'))

model.compile(loss='mse', optimizer='adam')
```

# Multi-layer NN

```
from tensorflow import keras
from tensorflow.keras.models import Sequential
from tensorflow.keras.layers import Dense

model = Sequential()

model.add(Dense(12, input_dim=8, activation='relu'))

model.add(Dense(12, activation='relu'))
model.add(Dense(12, activation='relu'))
model.add(Dense(12, activation='relu'))
model.add(Dense(12, activation='relu'))
model.add(Dense(12, activation='relu'))
model.add(Dense(12, activation='relu'))
model.add(Dense(12, activation='relu'))
model.add(Dense(12, activation='relu'))
model.add(Dense(12, activation='relu'))
model.add(Dense(12, activation='relu'))
model.add(Dense(12, activation='relu'))
model.add(Dense(12, activation='relu'))
model.add(Dense(12, activation='relu'))
model.add(Dense(12, activation='relu'))
model.add(Dense(12, activation='relu'))
model.add(Dense(12, activation='relu'))

model.add(Dense(1, activation='linear'))

model.compile(loss='mse', optimizer='adam')
```

# Activation functions: domain of $Y$ !

## 1. Bounded

- Heaviside (pre-2000)

$$\theta(x) \in [0, 1]$$

- sigmoid

$$\frac{1 + \tanh(x)}{2} \in [0, 1]$$

## 2. Unbounded

- ReLU

$$\max(0, x) = (x > 0) \cdot x \in [0, \infty]$$

- LeakyReLU ( $\alpha$  may be learned  $\rightarrow$  PReLU)

$$\max(0, x) + (x < 0) \cdot (\alpha x) \in [-\infty, +\infty]$$

- SELU

$$\lambda x \cdot (x < 0) + \lambda \alpha(e^x - 1) \cdot (x \geq 0) \in [-1.75, +\infty]$$

$\lambda, \alpha$  known

- Swish [link]

$$x \text{sigmoid}(x)$$

# How to train and predict a NN with Keras

```
from tensorflow import keras
from tensorflow.keras.models import Sequential
from tensorflow.keras.layers import Dense

P=20
K=1

model = Sequential()
model.add(Dense(12, input_dim=P activation='relu'))
model.add(Dense(12, activation='relu'))
model.add(Dense(12, activation='relu'))
model.add(Dense(K, activation='linear'))

model.compile(loss='mse', optimizer='adam')

# here, generate X, an N*P numpy array
#           and Y, the right answer
# X=...
# Y=...

model.fit(X,Y)

# X_test is a new input (numpy array of length N'*P)
# X_test=...

model.predict(X_test)
```

## Training in practice

- $N \gg 1$  samples: computing loss is expensive
- Batch size  $B < N$ : number of samples

$$\mathcal{L}(X) \simeq \mathcal{L}(X_{batch})$$

- for each batch
  - compute loss
  - backpropagation #stochastic gradients
- Number of times whole data set is used: epochs
  - for each epoch
    - for each batch
    - compute loss
    - backpropagation #stochastic gradients

# How to train and predict a NN with Keras

```
from tensorflow import keras
from keras.models import Sequential
from keras.layers import Dense

P=20
K=1

model = Sequential()
model.add(Dense(12, input_dim=P activation='relu'))
model.add(Dense(12, activation='relu'))
model.add(Dense(12, activation='relu'))
model.add(Dense(K, activation='linear'))

model.compile(loss='mse', optimizer='adam')

# here, generate X, an N*P numpy array
#           and Y, the right answer

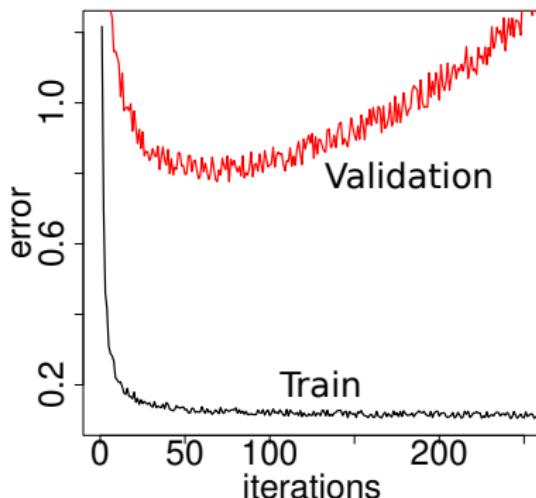
model.fit(X,Y, batch_size=32, epochs=10)

# X_test is a new input (numpy array of length N'*P)
# X_test=...

model.predict(X_test)
```

## How many epochs? batch size?

1. Split train data into train + validation
2. Early stopping when error on validation stop decreasing



## Model calibration: why ML?

- ML requires a stochastic model / candidate distribution
- Some parameters may be hard to estimate with ML
- Estimation is more generic
- Best estimator usually unknown

# Model calibration

- Stochastic model  $\mathcal{M}(\theta)$ ,  $\theta \in \mathbb{R}^K$

noise  $\rightarrow$  time series  $x = \{x_1, \dots, x_T\}$

- Calibration:

$$\hat{\theta}(x) = \hat{M}^{-1}(x)$$

ML/NN predicts  $\theta$  from model output  $\equiv \mathcal{M}^{-1}(X)$

# Model calibration: recipe

- choose  $J$  values of  $\theta_i$  (regularly spread)
- If stochastic model
  - for each  $\theta_i$ , generate  $M$  outputs

$$X^{(\theta_i)} = \begin{pmatrix} X_{1,1}^{(\theta_i)} & \cdots & X_{1,T}^{(\theta_i)} \\ \vdots & \vdots & \vdots \\ X_{M,1}^{(\theta_i)} & & X_{M,T}^{(\theta_i)} \end{pmatrix} \sim \begin{pmatrix} \theta_i \\ \vdots \\ \theta_i \end{pmatrix} = Y^{(\theta_i)}$$

- stack the matrices

$$X = \begin{pmatrix} X^{(\theta_1)} \\ \vdots \\ X^{(\theta_J)} \end{pmatrix} \in \mathbb{R}^{MJ \times T} \quad Y = \begin{pmatrix} Y^{(\theta_1)} \\ \vdots \\ Y^{(\theta_K)} \end{pmatrix} \in \mathbb{R}^{MJ \times K}$$

- Else  $M = 1$

# Reference papers in finance

[link]

## Model Calibration with Neural Networks

Andres Hernandez  
IBM Risk Analytics  
`andres.hernandez@gmx.net`

July 20, 2016

- Quantity  $Q = \mathcal{M}(\theta_{\text{model}}, \xi_{\text{external}})$
- Data:  $\{Q_{\text{market}}, \xi_{\text{external}}\}$
- Learn

$$\theta^* = \arg \min_{\theta} \text{Cost}(\theta, \{Q_{\text{market}}, \xi_{\text{external}}\}) = \text{fct}(\{Q_{\text{market}}, \xi_{\text{external}}\})$$

# Reference papers in finance

[link]

Vol. 21, No. 1, 11–27, <https://doi.org/10.1080/14697688.2020.1817974>



## Deep learning volatility: a deep neural network perspective on pricing and calibration in (rough) volatility models

BLANKA HORVATH<sup>\*†</sup>, AITOR MUGURUZA<sup>‡</sup> and MEHDI TOMAS<sup>§</sup>

<sup>\*</sup>King's College London, The Alan Turing Institute London, UK

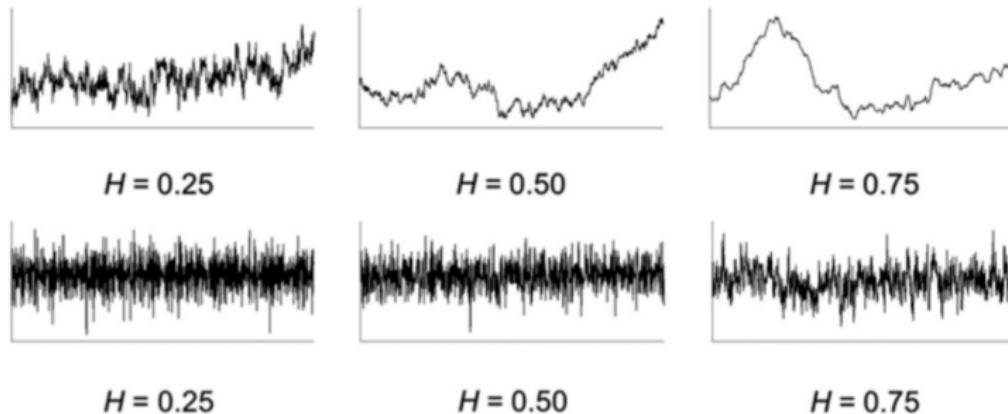
<sup>‡</sup>Imperial College London & NATIXIS London, UK

<sup>§</sup>CMAP & LadHyx, École Polytechnique, Palaiseau, France

When model  $\rightarrow Q$  is expensive (e.g. volatility map)

1. train  $NN_1$   $Q[\mathcal{M}(\theta, \xi_{external})] \rightarrow \tilde{Q}(\theta, \xi_{external})$  (surrogate)
2. train  $NN_2$  to solve calibration

## Case study: Hurst exponent $H$



**FIGURE 1 | Top row: example series of fractional Brownian motions (fBm) for three typical values of the scaling exponent.** The central graph represents an ordinary Brownian motion ( $H=0.5$ ). The left graph shows an anti-persistent fBm ( $H=0.25$ ) and the right graph a persistent fBm ( $H=0.75$ ). The corresponding fractional Gaussian noises series (fGn) are displayed in the bottom row.

Marmelat et al. (2012) [link]

## Standard diffusion

- $X_t = X_{t-1} + \epsilon_t$

- If  $\epsilon_t$  i.i.d.,

$$E[(X_t - X_0)^2] = cst \times t = cst \times t^{2H}$$

- In physicist's words: envelope growth

$$|X_t - X_0| \sim \sqrt{t} = t^H$$

$H$  : Hurst exponent

- Diffusion / Wiener process:

$$H = 1/2$$

## How to generate anomalous diffusion

fractional Brownian motion (fBm)  
(Mandelbrot and Van Ness (1968) [link])

- Notation: position  $W_t$
- $W_t - W_s$  is stationary

$$P(W_t - W_s) = P(W_{t-s})$$

- Hyp: Gaussian increments with

$$E(dW_t) = 0, \quad E(dW_t^2) = 1$$

## fMB: properties

- Covariance

$$E(W_t W_s) = \frac{1}{2} \left( |t|^{2H} + |s|^{2H} - |t-s|^{2H} \right) \begin{cases} > 0 & H > 1/2 \\ \min(t, s) & H = 1/2 \\ < 0 & H < 1/2 \end{cases}$$

- Self-affine time

$$t \rightarrow \lambda t$$

$$W_{\lambda t} \xrightarrow{\text{law}} \lambda^H W_t$$

- Integral representation

$$\begin{aligned} W_t = W_0 + \frac{1}{\Gamma(H+1/2)} &\left\{ \int_{-\infty}^0 \left[ (t-s)^{H-1/2} - (-s)^{H-1/2} \right] dW_s \right. \\ &\left. + \int_0^t (t-s)^{H-1/2} dW_s \right\} \end{aligned}$$

## Simulation of fBm

1. Straightforward integration from Gaussian variables
2. Given that the covariance between  $B_s$  and  $B_t$  is known, use Cholesky decomposition.
3. ...

## How to measure $H$ ? Scale plot

Since

$$E[(W_t - W_0)^2] \propto t^{2H}$$

Fit

$$\log E[(W_t - W_0)^2] = a + 2H \log t$$

# Rescaled Range method (Hurst)

## Range method

- Detrend data

$$X_t = W_t - \hat{\mu}, \quad t = 1, \dots, T$$

- Compute  $S_t$  : standard deviation of  $X_t$
- Compute running minimum and maximum

$$m_t = \min(X_1, \dots, X_t)$$

$$M_t = \max(X_1, \dots, X_t)$$

→ NOT robust!

- Range

$$R_t = M_t - m_t$$

- Rescaled range

$$R_t / S_t \simeq H$$

## Fourier Method

- Fourier transform: from real signal to frequencies → spectrum
- Non-local kernel  $e^{i\pi kx}$
- Flandrin (1989) [link]: for fBm, average spectrum

$$S(\omega) \propto \frac{1}{|\omega|^{2H+1}}$$

- Compute spectrum, fit with power-law →  $H$

# Wavelets

- Local kernels  $\phi$

$$\int \Phi(t) dt = 0, \quad \int \Phi(t)^2 dt = 1$$

Orthogonality to low-order polynomials

$$\int t^m \Phi(t) dt = 0 \quad 0 \leq m \leq n$$

- Example: Morlet (mother) wavelet

$$\Phi(t) = \frac{c}{\pi^{1/4}} e^{-\frac{t^2}{2}} \left( e^{it} - \kappa \right) \sim e^{-\frac{t^2}{2}} e^{it}$$

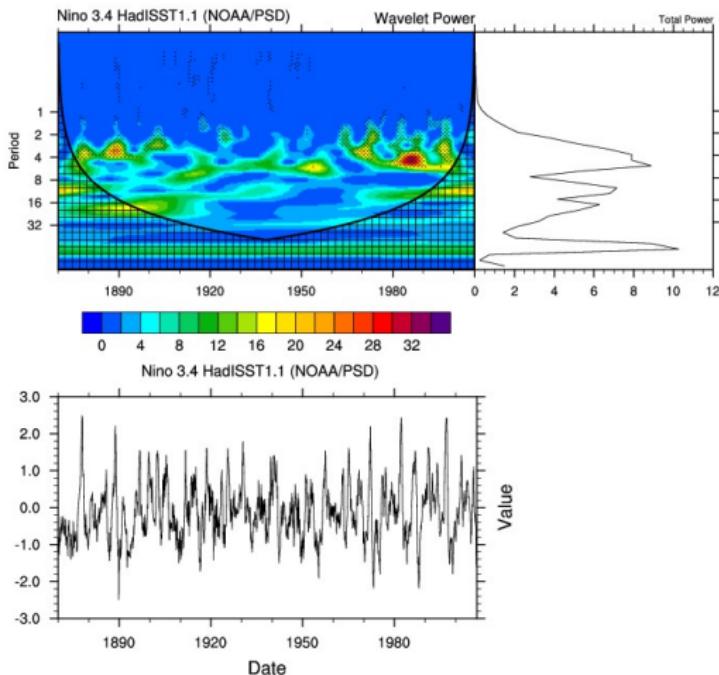
- Vary scale and time shift  $\rightarrow$  wavelet basis

$$\Psi_{a,b}(t) = \Phi \left( \frac{t-b}{a} \right)$$

# Wavelet transform

Cf Fourier

$$\mathcal{W}[h](a, b) = \frac{1}{\sqrt{a}} \int \Psi_{a,b}^*(t) h(t) dt$$



[link]

## Estimation of $H$ with wavelets

Simonsen *et al.* (1998) [link]

Reminder: if  $h$  is self-affine,  $h(\lambda t) = \lambda^H h(t)$

$$h(t) = \lambda^{-H} h(\lambda t)$$

thus

$$\begin{aligned}\mathcal{W}[h(t)](a, b) &= \mathcal{W}[\lambda^{-H} h(\lambda t)](a, b) \\ &= \dots \\ &= \lambda^{-H-1/2} \mathcal{W}[h(t)](\lambda a, \lambda b)\end{aligned}$$

thus finally

$$\mathcal{W}[h(t)](\lambda a, \lambda b) = \lambda^{H+1/2} \mathcal{W}[h(t)](a, b)$$

## Estimation of $H$ with wavelets

Exploit

$$\mathcal{W}[h(t)](\lambda a, \lambda b) = \lambda^{H+1/2} \mathcal{W}[h(t)](a, b)$$

1. Average time shift out: define

$$W[h](a) = \text{mean}(|\mathcal{W}[h](a, b)|)_b$$

2. Scaling: take several  $\lambda$ s

$$W[h](\lambda a) \simeq \lambda^{H+1/2} W[h](a)$$

3. Fit

$$\log W[h](\lambda a) = (H + 1/2) \log \lambda W[h](a) + cst$$

$\rightarrow$  slope  $H + 1/2$

## Other methods

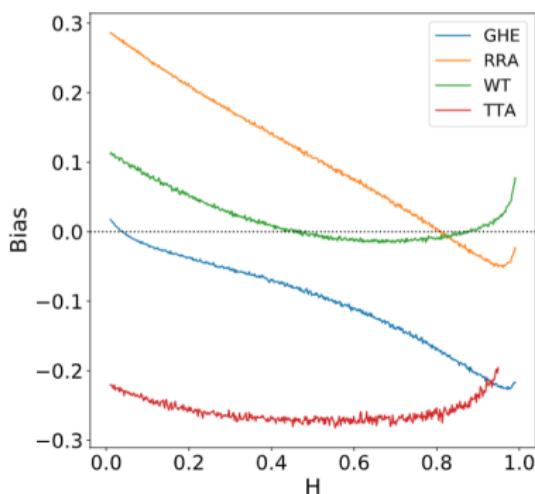
- Detrended fluctuation analysis (large biases)
- Total Triangle area (2020) [link]
- Distance with self-affine pdfs.
- Ratio of scaling [Garcin (2021)]

$$\hat{H}_t = \frac{1}{2 \log(\tau_1/\tau_2)} \log \left( \frac{(T - \tau_2) \sum_{i=0}^{T-\tau_1} (p_{t-i} - p_{t-i-\tau_1})^2}{(T - \tau_1) \sum_{i=0}^{T-\tau_2} (p_{t-i} - p_{t-i-\tau_2})^2} \right)$$

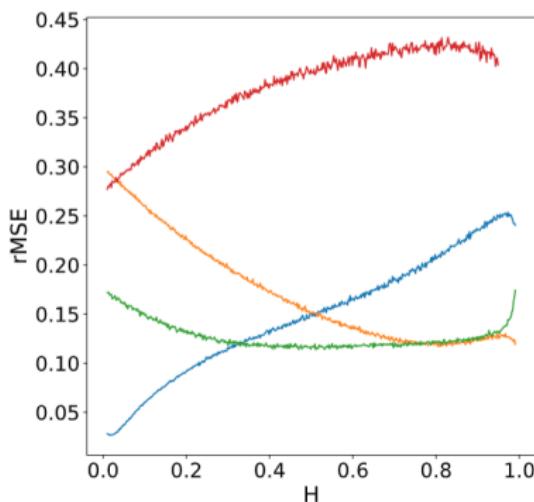
# Problem: short time series

Nguyen and Challet (2021)  $T = 100$ .

GHE=moment method, RRA=rescaled range, WT=wavelets,  
TTA= total triangle area



(a) Bias



(b)  $rMSE$

# Estimate $H$ with DL

MultiLayerPerceptron (Dense) layers:

Rules of thumb

- # neurons first layer = dim\_input
- Layers structure
  - Rectangle □▷
  - Diamond ◄▷
  - Triangle ▷
  - Bottleneck ▷ ◄▷
- Activation functions: some intuition from first principles

## Estimate H with DL: convolutional networks

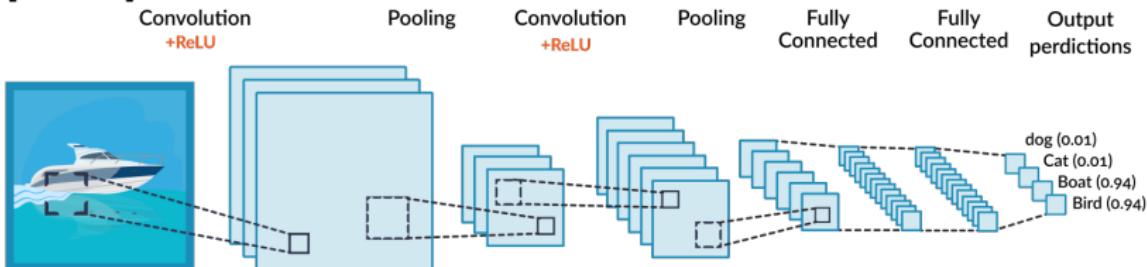
- fBm: convolution
- Generically, convolution between  $\{x_i\}$  and kernel  $\phi$

$$\sum_{k=-K/2}^{K/2} \phi_k x_{i-k} \quad i = 1, \dots, T$$

- $K$  : kernel size

# CNNs for images: typical architecture

[source]



- Kernels in parallel
- Smaller images, more kernels
- One feature per kernel

# CNNs for images

[source]



- Kernels in parallel
- One feature per kernel
- The deeper, the more and more complex the features
- Interactive visualization [[link](#)]

# CNNs

- Convolution between  $\{x_i\}$  and kernel  $\phi$  (learned)

$$\sum_{k=-K/2}^{K/2} \phi_k x_{i-k} \quad i = 1, \dots, T$$

- Then activation of linear convolution

$$\tilde{x}_i = \Phi \left( \sum_{k=-K/2}^{K/2} \phi_k x_{i-k} \right)$$

- From Keras help:

```
tf.keras.layers.Conv1D(  
    filters, kernel_size, strides=1, padding='valid',
```

- `filters` = number of parallel kernels
- `kernel_size` =  $K$  (cf above)
- `padding='same'` keeps dimension of input by adding zeros at the start and end of  $x$
- `padding='causal'` adds zero at the start (see [link]) and only takes input from the past

## CNN (1D): other ingredients

### Between layers

- pooling: reduce dimension, filter
  - MaxPooling:  $\max_{j \in [i-\delta, i+\delta]} (\tilde{x}_j)$ ;  $\delta$ : parameter
  - AveragePooling
- dropout: random neurons during training, avoids overfitting
  - rate = probability to disable a given neuron during training
  - rules of thumb: 0.5 in hidden layers, 0.8 in input layers
  - only during `.fit()`, not `.predict()`
  - equivalent to weight constraints

## $H$ from CNNs

Stone, Quant. Fin. (2020) [link] CNNs

- the first layer, with 32 kernels;
- max pooling layer;
- a dropout layer, with rate = 0.25;
- the second layer, with 64 kernels;
- max pooling layer;
- a dropout layer, with rate = 0.25;
- the third layer, with 128 kernels;
- max pooling layer;
- a dropout layer, with rate = 0.4;
- a dense layer with 128 units;
- a dropout layer, with rate = 0.3.

Leaky ReLU ( $\alpha = 0.1$ ), maxpooling of size 3, 30 epochs, batch size=64

## Challenge: image to $H$

- Time series are images
- Input a time series as its image
- Use Conv2D networks

# CNN $\longleftrightarrow$ wavelets

- CNN

$$\Phi \left( \sum_{k=-K/2}^{K/2} \phi_k x_{i-k} \right)$$

and wavelets are convolutions

- Wavelets fix a mother function
- **Maths:** S. Mallat 'CNNs really are wavelets' [link]  
→ pre-filtering with wavelets helps learning much faster
- **DL:** what is a wavelet? Who cares?

## Bonus topic: early stopping

- Until now, fix number of epochs
- Idea: stop when learning is useless
- Early stopping [link]: stop when loss or val\_loss does not decrease
  - minimum/average over the last  $time_{patience}$  steps
  - either train or validation loss
- code

```
from tensorflow.keras.callbacks import EarlyStopping  
es = EarlyStopping(monitor='val_loss', mode='min', patience=10)  
model.fit(      , callbacks=[es])
```