

Deep Learning in Finance

option pricing

Damien Challet
damien.challet@centralesupelec.fr

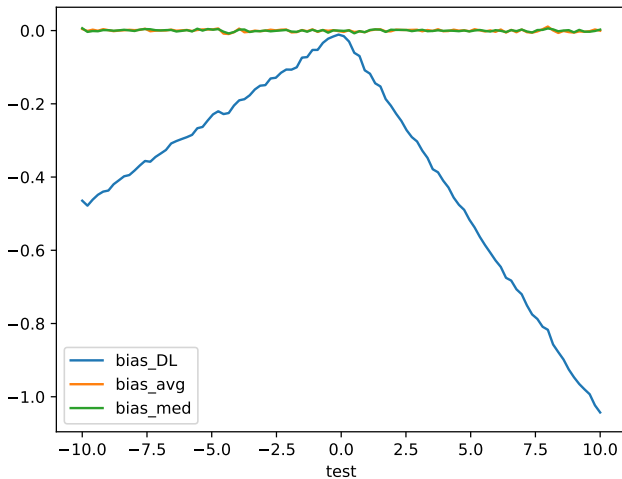
29th October 2024

Predictors: NNs vs trees

1. Tree-based methods cannot add, multiply, compare and lag columns
 - Additional predictors must be added
 - NNs can do that (learning to multiply is harder)
2. Tree-based methods cannot extrapolate beyond previously seen values
 - NNs can do that (not very well)

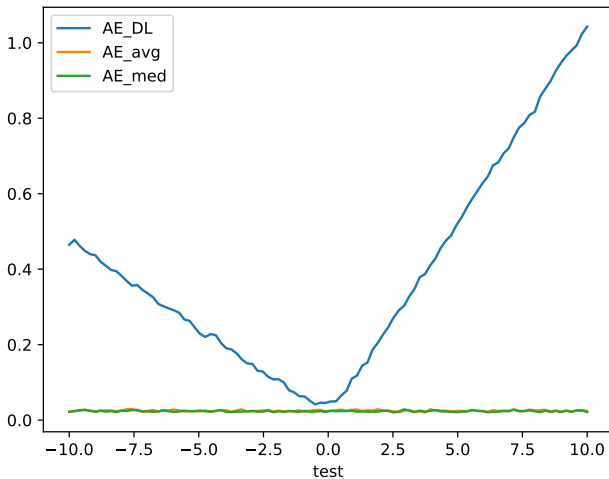
DL and sample mean: bias

DL: train for $\mu \in [-2, +2]$



DL and sample mean: variance

Mean absolute error: mainly from bias



Hard problems in Finance

1. Modelling and calibration
2. Solving high-dimensional equations
3. Prediction
4. Portfolio construction

Options


Insurance / bet: price/volatility change

option \equiv right to buy (*call*) or to sell (*put*) an asset at price K (strike).

- European option: single exercise time T (maturity)
- American option: any time until expiry
- barrier options
- exotic options

Payoff shape

Strike $K = 100$



from [link]

Options variables

- asset price p_t
- strike price K
- time to maturity T
- interest rate r

→ option price at time $t = 0$?

European option price

- Hyp: call option, strike K , time to maturity T
- Option price C = fair price

- Naively,

$$C = E[\max(p_T - K, 0)]$$

- Interest rate?

C can be invested in a risk-free asset

$$\begin{aligned} C(1+r)^T &= E[\max(p_T - K, 0)] \\ &= \int_K^{\infty} P(p_T)(p_T - K)dp_T \end{aligned}$$

→ model of price needed

Bachelier's model

- Price evolution

$$p_{t+1} = p_t + \eta_t$$
$$\eta_t \sim \mathcal{N}(0, \sigma_B^2)$$

- Normal diffusion

$$P(p_T | p_0) = \frac{1}{\sqrt{2\pi T} \sigma_B} e^{-\frac{(p_T - p_0)^2}{2T\sigma_B^2}}$$

- Additive world

Black-Scholes model

- Price evolution

$$\begin{aligned}\log p_{t+1} &= \log p_t + \eta_t \\ \eta_t &\sim \mathcal{N}(0, \sigma_{BS}^2)\end{aligned}$$

- Log-normal diffusion

$$P(\log p_T | \log p_0) = \frac{1}{\sqrt{2\pi T}\sigma} e^{-\frac{(\log p_T - \log p_0)^2}{2\sigma^2 T}}$$

- N.B.: no jumps, constant volatility

Fair price in discrete time with hedging

- $t \in \{0, \dots, T\}$
- v_t : number of shares invested in underlying asset

$$C(1+r)^T = E[\max(K - p_T, 0)] \\ - \sum_{t=0}^{t-1} E \left[(1+r)^{T-k-1} v_t (p_{t+1} - p_t - r p_t) \right]$$

- v_t : hedging (reduce or eliminate risk)

Optimal hedging

- Capital evolution of option underwriter:

$$\begin{aligned}\Delta W = & C(1+r)^T - [\max(K - p_T, 0)] \\ & + \sum_{t=0}^{t-1} \left[(1+r)^{T-k-1} v_t (p_{t+1} - p_t - rp_t) \right]\end{aligned}$$

- Residual risk = $\min_{\{v_t\}} \text{var}(\Delta W)$
- Continuous time & Gaussian increments

→ Black-Scholes miracle:

zero residual risk with optimal hedging

Black-Scholes option price

Continuous-time limit: $T = N\tau$, $\tau, 1/N \rightarrow 0$,

call price

$$C_{BS}(p_0, K, T, r, \sigma) = p_0 P_{>}[y_-] - Ke^{-rT} P_{>}[y_+]$$

$$P_{>}(x) = \frac{1 + \operatorname{erf}(x/\sqrt{2})}{2}$$

$$y_{\pm} = \frac{\log(p_0/K \cdot e^{-rT}) \pm \sigma^2 T/2}{\sigma\sqrt{T}}$$

Price scaling

$$C_{BS}(p_0, K, T, r, \sigma) = p_0 P_{>}[y_-] - Ke^{-rT} P_{>}[y_+]$$


$$P_{>}(x) = \frac{1 + \operatorname{erf}(x/\sqrt{2})}{2}$$
$$y_{\pm} = \frac{\log(p_0/Ke^{-rT}) \pm \sigma^2 T/2}{\sigma\sqrt{T}}$$

Setting $k = \frac{Ke^{-rT}}{p_0}$

$$\frac{C_{BS}(k, T, r, \sigma)}{p_0} = P_{>}[y_-] - kP_{>}[y_+]$$

Payoff shape

Strike $K = 100$



from [link]

Black-Scholes option price

Continuous-time limit: $T = N\tau$, $\tau, 1/N \rightarrow 0$,

Put price

$$P_{BS}(p_0, K, T, r, \sigma) = -p_0 P_{>}[-y_-] + Ke^{-rT} P_{>}[-y_+]$$

$$P_{>}(x) = \frac{1 + \operatorname{erf}(x/\sqrt{2})}{2}$$

$$y_{\pm} = \frac{\log(p_0/K \cdot e^{-rT}) \pm \sigma^2 T/2}{\sigma\sqrt{T}}$$

Parameters of Black and Scholes

- Known: p_0, K, T
- Possibly known: r
- Optimised: hedging $\{v_t\}$
- **Estimated:** σ
- **Assumed:** continuous time, Gaussian returns

Black-Scholes vs reality: returns

Heavy-tailed skewed price returns+stochastic volatility

→ implicit volatility:

$$\underbrace{C_{BS}(p_0, K, T, r, \sigma_{impl})}_{BS \text{ world}} = \underbrace{\text{option price}(p_0, K, T, r)}_{\text{real world}}$$

- Heavy tails → smile
- Skewed returns, leverage effect → skewed smile

Non-constant volatility

→ stochastic volatility model

→ non-explicit option price formulae (e.g. Heston, etc.)

N.B.: additional risk: volatility of volatility

Deep learning and option pricing

- | | |
|-------------------------|----------------------|
| 1. Model-driven | speed up calibration |
| 2. Data-driven approach | Ito who? |
| 3. Model + Data | improve on 1. and 2. |

1. Deep learning option prices from model

1. NN fitted to reproduce expensive model output
 - 1.1 Monte-Carlo
 - 1.2 Heston
 - 1.3 Rough Bergomi [\[link\]](#)
2. Find best parameters of the NN-fitted model
 - 2.1 with standard optimisation
 - 2.2 add no-arbitrage conditions
 - 2.3 NN [\[link\]](#)

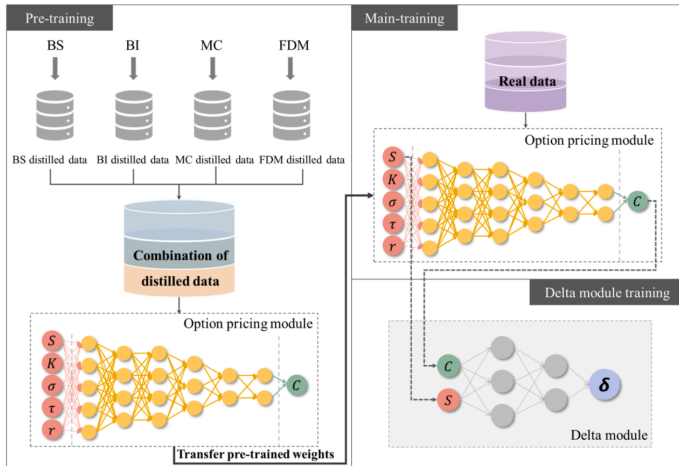
2. Data-driven option price

1. Fixed parameters (K, p_0, r, T)
2. Add estimated parameters σ , skew, tail exponent
3. Fit option prices with DL \rightarrow option pricer

3. Deep learning option prices: prices from model and data

Example: Jang *et al.* (2020)

[<https://doi.org/10.1016/j.inffus.2020.12.010>]



3. Deep learning option prices from model

Improvement:

1. Give what you know as inputs
2. Learn residuals

Example: Funahashi (2021) [\[link\]](#)

- Learn $C_3 - C_{data}$, where $C_3 =$ third order approximation of C_{BS}
 - Smaller training time
 - More robust and stable
 - Faster calculations

Solving complex option pricing equations

Wang et al (2022) [link]

- Physics informed NN (PINNS) (2019) [link] (11300 citations)

To describe the Physics-Informed Neural Networks for solving PDEs, let us consider a general PDE with the solution $u(\mathbf{x})$ defined on a domain $\Omega \subset \mathcal{R}^d$:

$$f\left(\mathbf{x}; \frac{\partial u}{\partial x_1}, \dots, \frac{\partial u}{\partial x_d}; \frac{\partial^2 u}{\partial x_1 \partial x_1}, \dots, \frac{\partial^2 u}{\partial x_1 \partial x_d}; \dots\right) = 0, \quad \mathbf{x} := (x_1, \dots, x_d)' \in \Omega, \quad (2)$$

supplemented with proper boundary conditions

$$\mathcal{B}(u, \mathbf{x}) = 0 \quad \text{on} \quad \partial\Omega, \quad (3)$$

where $\mathcal{B}(u, \mathbf{x})$ can be Dirichlet, Neumann, Robin, or even periodic boundary con-

Solving complex option pricing equations

Physics informed NN (PINNS) (2019) [link]

The major steps of the PINNs algorithm for solving a PDE are as follows (cf. Lu et al. 2021):

Step 1 Construct a neural network $\hat{u}(\mathbf{x}; \theta)$ with parameters $\theta = \{W_l, \mathbf{b}_l\}_{1 \leq l \leq L}$.

Step 2 Specify the training data \mathcal{T}_{pde} and \mathcal{T}_{ib} of sizes N_{pde} and N_{ib} , respectively, to satisfy the physics imposed by the PDE and the boundary/initial conditions. Here $\mathcal{T}_{pde} \subset \Omega$ and $\mathcal{T}_{ib} \subset \partial\Omega$ are points inside the domain Ω and on the domain boundary $\partial\Omega$, respectively.

Step 3 Specify a loss function defined as a weighted sum of the L^2 errors for the PDE and the boundary/initial conditions:

$$\begin{aligned} \mathcal{L}(\theta) := & \omega_{pde} \cdot \frac{1}{N_{pde}} \sum_{\mathbf{x} \in \mathcal{T}_{pde}} \left\| f\left(\mathbf{x}; \frac{\partial \hat{u}}{\partial x_1}, \dots, \frac{\partial \hat{u}}{\partial x_d}; \frac{\partial^2 \hat{u}}{\partial x_1 \partial x_1}, \dots, \frac{\partial^2 \hat{u}}{\partial x_1 \partial x_d}; \dots\right) \right\|_{L^2(\Omega)}^2 \\ & + \omega_{ib} \cdot \frac{1}{N_{ib}} \sum_{\mathbf{x} \in \mathcal{T}_{ib}} \left\| \mathcal{B}(\hat{u}; \mathbf{x}) \right\|_{L^2(\partial\Omega)}^2, \end{aligned} \quad (4)$$

where ω_{pde} and ω_{ib} are the weights. Those derivatives in the loss function (4) are evaluated by the automatic differentiation (AD) provided by Tensorflow (Abadi

Solving complex option pricing equations

Wang et al (2022) [link]

Non-linear BS equation (with transaction costs)

Let us consider the nonlinear Black–Scholes equation (Lesmana et al., 2013):

$$\frac{\partial V}{\partial \tau} - \frac{1}{2} \sigma^2(\tau, S, \frac{\partial V}{\partial S}, \frac{\partial^2 V}{\partial S^2}) S^2 \frac{\partial^2 V}{\partial S^2} - rS \frac{\partial V}{\partial S} + rV = 0, \text{ on } (0, \infty) \times (0, T], \quad (10)$$

where the volatility σ is a function of time τ , the underlying stock price S , the Delta $\frac{\partial V}{\partial S}$, and the Gamma $\frac{\partial^2 V}{\partial S^2}$.

Two-asset case

$$\begin{aligned} \frac{\partial V}{\partial \tau} - \frac{\sigma_1^2 S_1^2}{2} \frac{\partial^2 V}{\partial S_1^2} - (r - \delta_1) S_1 \frac{\partial V}{\partial S_1} + rV \\ - \frac{\sigma_2^2 S_2^2}{2} \frac{\partial^2 V}{\partial S_2^2} - (r - \delta_2) S_2 \frac{\partial V}{\partial S_2} - \rho \sigma_1 \sigma_2 S_1 S_2 \frac{\partial^2 V}{\partial S_1 \partial S_2} = 0, \quad \text{on } (0, \infty)^2 \times (0, T], \end{aligned} \quad (16)$$

where the option value $V(S_1, S_2, \tau)$ depends on the asset prices S_i ($i = 1, 2$) and the time to maturity τ . The parameter r is the interest rate, σ_i and δ_i ($i = 1, 2$) denote the volatilities and the dividend rates for the two assets, respectively, and ρ represents the correlation coefficient between S_1 and S_2 .

Solving complex option pricing equations


Wang et al (2022) [link]

Two-asset case

$$\begin{aligned} \frac{\partial V}{\partial \tau} - \frac{\sigma_1^2 S_1^2}{2} \frac{\partial^2 V}{\partial S_1^2} - (r - \delta_1) S_1 \frac{\partial V}{\partial S_1} + rV \\ - \frac{\sigma_2^2 S_2^2}{2} \frac{\partial^2 V}{\partial S_2^2} - (r - \delta_2) S_2 \frac{\partial V}{\partial S_2} - \rho \sigma_1 \sigma_2 S_1 S_2 \frac{\partial^2 V}{\partial S_1 \partial S_2} = 0, \quad \text{on } (0, \infty)^2 \times (0, T], \end{aligned} \quad (16)$$

where the option value $V(S_1, S_2, \tau)$ depends on the asset prices S_i ($i = 1, 2$) and the time to maturity τ . The parameter r is the interest rate, σ_i and δ_i ($i = 1, 2$) denote the volatilities and the dividend rates for the two assets, respectively, and ρ represents the correlation coefficient between S_1 and S_2 .

Fig. 5 Example 3. Plots of the numerical solution at the final time T (Top); the training history (Bottom)



Deep learning: regression + no-arbitrage constraints

Ackerer, Takasovska and Vatter (2020),
Deep Smoothing of the Implied Vol Surface [link]

- Implied volatility ansatz (residuals)

$$\sigma = \sigma_{NN} \sigma_{prior}$$
$$\omega(k, \tau) = \sigma^2 \tau$$

where σ_{prior} comes from a model and τ : time to maturity

- 6 shape and no-arbitrage constraints

DL: regression + constraints

Proposition 2.2 Roper [48, Theorem 2.9] Let $S > 0$, $r = \delta = 0$, and $\omega : \mathbb{R} \times [0, \infty) \mapsto \mathbb{R}$. Let ω satisfy the following conditions:

C1) (Positivity) for every $k \in \mathbb{R}$ and $\tau > 0$, $\omega(k, \tau) > 0$.

C2) (Value at maturity) for every $k \in \mathbb{R}$, $\omega(k, 0) = 0$.

C3) (Smoothness) for every $\tau > 0$, $\omega(\cdot, \tau)$ is twice differentiable.

C4) (Monotonicity in τ) for every $k \in \mathbb{R}$, $\omega(k, \cdot)$ is non-decreasing, $\ell_{\text{cal}}(k, \tau) = \partial_{\tau}\omega(k, \tau) \geq 0$, where we have written ∂_{τ} for $\partial/\partial\tau$.

C5) (Durrleman's Condition) for every $\tau > 0$ and $k \in \mathbb{R}$,

$$\ell_{\text{but}}(k, \tau) = \left(1 - \frac{k \partial_k \omega(k, \tau)}{2\omega(k, \tau)}\right)^2 - \frac{\partial_k \omega(k, \tau)}{4} \left(\frac{1}{\omega(k, \tau)} + \frac{1}{4}\right) + \frac{\partial_{kk}^2 \omega(k, \tau)}{2} \geq 0,$$

where we have written ∂_k for $\partial/\partial k$ and ∂_{kk} for $\partial^2/(\partial k \partial k)$

C6) (Large moneyneess behaviour) for every $\tau > 0$, $\sigma^2(k, \tau)$ is linear for $k \rightarrow \pm\infty$.


C1, C2, C3 are guaranteed by activation function

Constraints help generalization abilities

DL: regression + constraints

Loss function. We fit the network parameters and prior parameters θ by minimizing the loss function

$$\mathcal{L}(\theta) = \mathcal{L}_0(\theta) + \sum_{j=1}^6 \lambda_j \mathcal{L}_{C_j}(\theta) \quad (6)$$



Nonarbitrage constraint: put-call parity

- C : call price
 - P : put price
 - S_t : spot (current) price
 - K : strike price
 - T : time to maturity
 - r : interest rate

$$C - P = S - e^{-rT}K$$

Custom Loss in Keras

- Usual losses [doc]

```
model.compile(...,loss='mse') or 'mae' or  
...
```

- Generically [doc][tf math]

```
import tensorflow as tf
```

```
def myloss(y_true, y_pred):  
    diff2 = tf.math.square(y_true - y_pred)  
    return tf.math.reduce_mean(diff2, axis=-1)
```

```
model.compile(loss=myloss)
```

(advanced): gradients in custom loss

- `myloss(y_pred, y_true)`
- custom loss with gradients: use `GradientTape` [doc]
[stackoverflow]
- other ways to impose hard/soft constraints: Donti, Rolnik, Zolter (2021) [link]

Early stopping

- Until now, fix number of epochs
- Idea: stop when learning is useless
- Early stopping [link]: stop when loss or val_loss does not decrease
 - minimum/average over the last $time_{patience}$ steps
 - either train or validation loss

- code

```
from tensorflow.keras.callbacks import EarlyStopping
es = EarlyStopping(monitor='val_loss', mode='min', patience=10)
model.fit(          , validation_split=0.3, callbacks=[es])
```

