

Deep Learning in Finance

GANs II

Damien Challet
damien.challet@centralesupelec.fr

12th November 2024

Lecture plan

1. More advanced GANs
2. Methodological issues and improvements


Naive GANs as market simulators

- Input noise $\eta \in \mathbb{R}^{N_F=d \times T}$ with $\eta_i \sim \mathcal{N}(0, 1)$
- Output price returns $r_t \in \mathbb{R}^T$
- Stylized facts?
 - heavy tailed returns hard/impossible
 - long memory of volatility hard/impossible
- better than i.i.d. r_t ?

discriminator says yes

Results: r_t vs t , Gaussian features

GAN: MLP, $T = 128$, $N_F = 10$, ReLU



Theorems

[link] and [link]


1. CNNs cannot transform Gaussian noise into heavy-tailed noise
2. MLPs cannot transform Gaussian noise into heavy-tailed noise

Solutions:

1. Use another NN architecture
2. Use another type of noise
3. Transform data


Results: r_t vs t , rescaled, Gaussian features

GAN: MLP, $T = 128$, $N_F = 10$, ReLU, individually rescaled




Results: r_t vs t , Student-t features

GAN: MLP, $T = 128$, $N_F = 10$, ReLU



Results: r_t vs t , rescaled, Student-t features


GAN: MLP, $T = 128$, $N_F = 10$, ReLU, individually rescaled



What did the generator learn?

- To modulate the volatility from features
see stochastic volatility models

$$r_t^{(b)} = \sigma^{(b)} \epsilon_t^{(b)}, \text{ for a given } \eta^{(b)}$$



- $P(\sigma_{sample})$ is heavy tailed
- ACF of σ : memory \equiv length of samples T

What did the discriminator learn?


Experiment

- increase scale of returns
- plot average \hat{p}_{true} vs scale

```
: 1 X_gauss=np.resize(np.random.normal(size=T*ntries),(ntries,T))
  2 |scales=np.linspace(0.5,10,20)
  3 y_avg=[]
  4 for myscale in scales:
  5     X_gauss_scaled=X_gauss*myscale
  6     y_gauss_scaled=discriminator.predict(X_gauss_scaled).flatten()
  7
  8     y_gauss_scaled_avg=np.mean(y_gauss_scaled)
  9     y_avg.append(y_gauss_scaled_avg)
10
11 plt.plot(scales,np.array(y_avg))
```

What did the discriminator learn?

SCALE!



If $P(\sigma_{sample}) \simeq P(\sigma_{real})$, GAN converges (weaker conditions apply)

Return timeseries generation from models

1. Long memory: fractional Brownian motion
→ Wiener process+power-law kernel
2. Heavy tails:
 - 2.1 external events → heavy tails features
 - 2.2 herding: internal state of agents/market

$$\sigma_t = F(\text{past}) \rightarrow \text{long memory} + \text{heavy tails}$$

- Better NN architecture
 - CNN with longer memory: Temporal CN (next slide)
 - stateful neural networks: GRU, LSTMs (soon)

GANs with TCNs

Wiese et al. (2019) [preprint]

- $T = 100$
- TCN: temporal convolutional networks \equiv Wavenets (Oord et al. 2016 [preprint])
stacked CNN with dilation ($\rightarrow O(\log N)$ coefficients)




Figure 6: Vanilla TCN with 4 hidden layers, kernel size $K = 2$ and dilation factor $D = 2$ (cf. van den Oord et al. (2016)).

[keras-tcn]

GANs a scenario generators

Mitigating overfitting [link]

Improving robustness of backtesting [link]

- Only touch ONCE a given set of data when backtesting
- Tweaking a strategy and testing it again \equiv data snooping
- Solutions
 1. Use statistical tools to account for multiple hypothesis testing
 2. Use another dataset

Multiple Hypothesis Testing

- For every strategy i , $H_0^{(i)} : E(\text{perf}_i) = 0$
- Control the fraction of false positive detection rate.
- Control the number of wrongly selected strategies.
- **Problem:** false negative rate can be very high.

Use another dataset


Naively:

- Aim: use a strategy on US equities
- Keep 20% of equities out of sample
- Calibrate on 80%
- **Problem:** large dependence between equities
- **Problem:** large dependence between countries

Robust backtest with a single asset

VIX prediction problem

Aim: predict when $VIX > 15$



Usual procedure



1. Sliding window
2. Tweak moving average of log VIX or stochastic volatility model

Prediction problem with Generative Model



1. Generate many VIX-like timeseries from training data set
2. Train a predictor of peaks on synthetic data
3. Test on real data

VIX > 15 prediction without data augmentation

From De Meer (2019) [link]



VIX>15 prediction with data augmentation



1. Open questions:

1.1 Why does it improve testing performance?

1.2 In what respect is it not overfitting of some kind?


[code]

conditional GANs


- Generate objects conditional on signal $s \in \mathbb{R}^S$
- Generator: input (η, s)
 - Discriminator: input $(G(\eta|s), s) \rightarrow P_{true}(y|s)$
- Example: scenario [link][link]
 - $s_t = (r_t, r_{t-1}, \dots, r_{t-S})$
 - Generator: outputs r_{t+1}, \dots, r_{t+k}
 - Discriminator: realistic or not

conGAN: example


Koshiyama *et al.* (2020) [link]



(d)



(e)



(f)

What to do with scenarios generators

Strategy \equiv model M_θ ; performance on data r : $\Pi[M_\theta(r)]$

1. Robust strategy parameter selection

- Best strategy on a single time series \equiv noise fitting
- Best strategy on many lookalike time series
 $\arg \max_\theta E(\Pi[M_\theta(r)])$

2. Aggregation of many overfitted strategies

- Best strategy for each scenario
- Overall strategy = average strategy
- Combination of weak learners \rightarrow strong learner (see trees)

Aggregations of weak strategies

Koshiyama *et al.* (2020) [link]




(g)



(h)


Methodological problems: convergence

Takayashi *et al.* (2019) [paper]



Convergence problems

MLP use_bias=False for both G and D .



How to improve convergence

Convergence problems \equiv game problems \equiv design problems

- Adapt G and D architecture
e.g. MLP + use_bias=False for both G and D
- Noisy labelling (Salimans *et al.* 2016 [paper]):
randomize a fraction of correct / false labels
- Better loss: compare full distributions of \hat{p}_i
WassersteinGAN [link]

How to (almost) solve convergence problems

Relativistic GANs (Jolicoeur-Martineau 2018 [preprint]):

'[...] the discriminator estimates the probability that the given real data is more realistic than fake data, on average'


$$L_D^{RaSGAN} = -\mathbb{E}_{x_r \sim \mathbb{P}} [\log (\bar{D}(x_r))] - \mathbb{E}_{x_f \sim \mathbb{Q}} [\log (1 - \bar{D}(x_f))] ,$$

where

$$\bar{D}(x) = \begin{cases} \text{sigmoid}(C(x) - \mathbb{E}_{x_f \sim \mathbb{Q}} C(x_f)) & \text{if } x \text{ is real} \\ \text{sigmoid}(C(x) - \mathbb{E}_{x_r \sim \mathbb{P}} C(x_r)) & \text{if } x \text{ is fake.} \end{cases}$$

[keras-relativistic-gan]

Example: relGANs for timeseries



Does convergence imply correct stylized facts?

Better generators: conditional GANs

Generator's input:

- feature vector (world, past returns, ...)
- noise vector




Figure 2.: An illustration of a conditional GAN pipeline.

[source]

- condGAN [preprint]
- ForGan [paper]
- Fin-GAN [preprint]

Better generators: loss

Ask what you want in generator loss

Fin-GAN: generator loss:

- Standard Generator loss
- + losses over prediction horizon
 - SR out-of-sample
 - hit ratio, ...

Better generators: timeseries representation

Pdf

- $r_t \rightarrow [0, 1]$: CDF: EV-GAN (Allouch *et al.* 2022 [link])
generate heavy tails consistent with data

Representation of whole timeseries

- Path signatures
 - Signature-GAN [preprint]
 - Sig-Wasserstein-GAN [preprint]
- Wavelet scattering spectra [link]

Perturbations to representations \rightarrow new scenarios

Better discriminator: what can a discriminator learn?

Without help, the discriminator easily learns

- moments (average, scale, etc)

Possibly

- Hurst exponent
- $\sum_{\tau} C_{|r|}(\tau)$
- Wavelets (CNNs)

Not

- quantiles, distribution tail exponent (see EV-GAN, Tail-GAN)

GAN for financial timeseries: a weird methodology

1. Train a GAN in an agnostic way
2. Convergence: happy
3. Look at timeseries.
4. Noise \rightarrow test a set of stylized facts

Unconsistent aims and methodology

Example: heavy tails treatment in GAN literature

[review 2021]





Figure 13: PDF - SAGAN


[Lemzi et al. 2021]



(b) Student's t -distribution with $\nu = 4$

Example: generating correlation matrices as images

Marti (2019) [preprint][link]



(Left) Empirical correlation matrix estimated on stocks returns; (Right) GAN-generated correlation matrix.

Nice looking, but

- generated matrices are not positive definite
- eigenvalue distributions not realistic

Be helpful, be explicit

- Pre-process data (remember: MLP cannot compute quantiles)
 - Hurst exponent of $|r_t|$
 - $\sum_{\tau=1} |C_r(\tau)|$: absence of linear autocorrelation
 - $\sum_{\tau>0} C_{|r|}(\tau)$: long memory
 - Non-Gaussian test (e.g. N-test [preprint])
 - Heavy-tail test: quantiles ratio
 - ...
- Post-process data: loss
 - Tail-GAN [preprint]: score-inspired loss for VaR/C-VaR

Discriminator with pre-processing abilities

Generalized moments method (GMM) [paper]

- Compute a set of K quantities $\{q_k\}$ from r_t (moments, ACF of moments, quantiles, etc)
- Implicit model calibration:
 - find model parameters θ so that $\{\hat{q}_k|\theta\} \simeq \{q_k^{\text{real data}}\}$
- Here: feed discriminator with $\{q_k|\theta\}$ that you kindly compute, plus possibly raw returns

Pre-processing: maths and stats

- math: \rightarrow `tf.math`
e.g. `np.sum()` \rightarrow `tf.math.reduce_sum()`
- stats: `import tensorflow_probability as tfp`

`tfp.stats.auto_correlation()`

How to help the discriminator: Lambda layers

- Python: lambda function
- Keras: Lambda layer: apply a function to the input of a layer
- Here: price returns \rightarrow discriminator \rightarrow statistics $\{a, b, c, \dots\}$
- Discriminator: first layer is a Lambda layer that calls the function computing the statistics
- Technical hint: convert list $[a,b,c]$ to tensor

```
return tf.transpose(tf.convert_to_tensor([a,b,c]))
```


Keras: functional models

- Several inputs, several sub-networks
- Merge inputs, outputs

```
|  
input_merged = Concatenate()([input_conds, input_eta])  
model=Dense(dim_conds,use_bias=use_bias,activation="PReLU")(input_merged)
```

- Shortcuts: merge input and model output (inception)

Technical note on model re-use: weights

```
#define model
model.compile(...)
model.fit(...)
model.saveweights('model_weights.h5')
```

Later, elsewhere

```
#define model
model.compile(...)
model.loadweights('model_weights.h5')    #instead of .fi
```

Technical note on model-reuse: architecture

- Each model can be described as a tree
- Model architecture saved as json [\[link\]](#)

Save

```
model_json = model.to_json()
with open("model.json", "w") as json_file:
    json_file.write(model_json)
model.save_weights('model_weights.h5')
```

Load

```
with open('model.json', 'r') as json_file:
    model_descr = json_file.read()
model = model_from_json(model_descr)
model.load_weights("model_weights.h5")
```